

Author Response Template

One section per cluster, four sentence roles per section

1 Is a template encouraged practice

Question 1 asks this directly, and records a suspicion alongside it: that venues are deliberately limiting what a response can contain, and that they may want only the factual problem corrected.

The suspicion about scope is right. The practice is encouraged, but for a document that is not this one.

Templating a response is ordinary and the filled-in artifacts are everywhere: community L^AT_EX templates, a Typst package, lab templates, a publisher’s fillable form. Every one of them is a journal response letter. They quote each reviewer comment, reply underneath it, cite the page and line where the manuscript changed, and compile to a document that travels beside a revised manuscript. No program committee publishes one. The strongest statement from a venue body is SIGPLAN’s guidance to its chairs, which encourages holding a response phase at all and says nothing about the shape of the response.

So the format is supplied for the wrong genre. A 750 word plain text box, one box for three reviewers, no revised manuscript to point into, and no line numbers, cannot carry a per-comment reply at all.

On the limiting, the call for papers of the IEEE Symposium on Security and Privacy for 2027 says it outright. The rebuttal “should exclusively focus on factual errors in the reviews and concrete questions posed by the reviewers,” and “New research results or other additional material cannot be provided.” The same page adds that failure to follow the instructions “will result in immediate rejection.” That is the scope, written by the people who set the limit. It is not a hint to be inferred from the word count.

What follows for this template. Scope is constrained and depth is not. The venues limit what a response may be about. None of them says how deeply the things in scope may be answered. A shallow answer therefore recovers nothing that a deep one would have cost.

2 Do people cluster as corrections and discussions

Question 3 asks whether that split is a recognized one and whether templates for it exist.

It is not recognized. DISAPERRE annotates 20,000 sentences from 506 review and response pairs at ICLR 2019 and 2020, and its scheme has fourteen response actions in three stances. Nothing in it splits corrections from discussions, and no venue’s guidance uses that division either. There is nothing to copy, so there is no cost to inventing a structure that suits the panel in front of you.

Three of its numbers are worth keeping. Answering a question is 32.76 percent of all response sentences, three times the next argumentative label. A response is mostly answers. Rejecting a criticism is 10.37 percent against 2.70 for conceding one, so refusing a reviewer, with a reason, is the common case rather than the risky one. Claiming a task is already done is 8.58 percent against 2.01 for promising it by the deadline, and that is at a venue with far more room to promise than this one.

One number cuts the other way. Contradicting something the review presented as fact is the rarest argumentative label in the set, 0.86 percent. The security venues invite that move ahead of every other. The thing most wanted here is the thing least done where the counting was possible.

So the clustering rule in this template is not corrections against discussions. It is root cause, which is the rule already written into the procedure on the site: the root of a comment is the thing that would have to change for the comment to go away, and comments sharing a root are answered once, as a group, with the group named.

3 The macros

Paste these into the response preamble. Every one of them is stripped by the exporter, so the working apparatus never reaches the submission form.

```
% One cluster. #1 the kind of engagement, #2 the subject.
\newcommand{\rcluster}[2]{\subsection*{#1: #2}}

% The four sentence roles, as reminders while drafting. Delete before export.
\newcommand{\ranswer}[1]{#1} % the direct answer, first
\newcommand{\rground}[1]{#1} % the fact, number, equation, or passage
\newcommand{\rconseq}[1]{#1} % what follows, formally and operationally
\newcommand{\rchange}[1]{#1} % what the revision does, checkable

% The reviewer's own phrase, echoed in the first clause. Marks it so the
% coverage tool can confirm every cluster carries one.
\newcommand{\recho}[1]{#1}

% A commitment that passed the gate. #1 configuration, #2 table it lands in.
\newcommand{\rcommit}[2]{#1, reported in #2}

% An offer, numbered, with what it would show. Never framed as required.
\newcommand{\roffer}[2]{#1, which would show #2}
```

4 The skeleton, one cluster

```
\rcluster{Kind}{Subject}
\recho{[the reviewer's own distinctive phrase]}: \ranswer{[the direct answer, as
asked]} \rground{[the fact, number, or passage that supports it]} \rconseq{[what it
means formally, and what it means for whoever would use the system]} \rchange{[what
the revision does, stated so a shepherd can check completion]}
```

Kind, chosen before drafting. A clarification of. A consideration of. A misunderstanding about. The hidden assumptions of. The motivation for. A title naming only a topic tells the reader nothing about what the section does with it.

The echo is the first clause, not a separate sentence. It costs about ten words per paragraph and it is the last thing to cut. Without it a correct sentence answers a question the reader has not

been reminded of.

The four roles are an order, not a length. Any of the four may be one clause. None may be absent, except that a cluster with nothing to change omits the change rather than inventing one.

5 Posture, selected by the classification

The classification step assigns every comment one of four kinds, and the kind decides the posture. This table is the mapping.

Kind	Premise	Posture
Misreading	The manuscript contradicts it	Observe that there appears to be a misunderstanding, then the correct fact. Never a statement that the reviewer is wrong. Treat it as evidence the text was unclear.
Defect	The manuscript confirms it	Agree first. Thank them where they spelled out the fix. Then the fix, without dwelling on how the error arose.
Scope question	Asks where the claim's boundary sits	Name the boundary. If the case sits outside the threat model, say so and say where the boundary is.
Request	Asks for work	Through the gate. Commit, decline with an honest reason, or offer.

The declining rule, both halves in one paragraph, always. Refute the phrasing where it is inaccurate, and honor the concern underneath it. A refusal without the second half reads as defensive. A concession without the first gives up ground the reviewer did not ask for.

6 The gate, before any commitment enters a draft

Three questions in order. No commitment passes without the author's sign-off.

1. Was the original choice justified, and is the justification written where the reviewer looked? If so, the answer is that justification and no work is promised.
2. Is the work feasible inside the response window or the shepherded revision? If not, decline with an honest reason.
3. Does it require advancement at the scale of a follow-up project? If so it is not promised, because promising work at that scale concedes the paper is a major revision at best.

Two constraints on what leaves the gate. A promise names a configuration and a table, because a shepherd can hold nothing else. A decline states a reason that is never effort and never the word limit.

Work you are unsure about becomes a numbered offer with what it would show, and the paper must not appear to depend on the reviewers accepting one.

7 Filling it

The order of the work is the nine step procedure on the website at `resources/research/respond/procedure.md`. It is not repeated here. The two steps that touch this file are step seven, which drafts one section per cluster, and step eight, which runs the wording pass and the per-question audit.

Only step seven writes. Everything before it extracts, reads, clusters, classifies, and gates.

Step seven expanded slot by slot, with the word budget, the refuse-to-start list and the exit tests, is `template_rebuttal_procedure.tex` beside this file.

8 The cut order, when it runs long

Three places, in this order, and none of them costs anything.

Restatement of the manuscript first. The reviewers have read the paper.

Spectator voice second. State your own numbers rather than reconstructing them from published figures in front of the reader.

Pointing instead of answering third. A section number is not an answer, since the reviewer already read that section, which is how they formed the comment.

If it still runs long, cut the section answering the least movable review.

9 What never reaches the form

Comment identifiers. Status markers. Approval labels. Notes on open decisions. Gates addressed to the author. Any working apparatus at all. The plain text is generated from this source rather than kept as a second copy by hand, because a second copy drifts and the drift is found at the form.

10 Sources

- The procedure, the drafting shape, the postures, the phrasing shelf, and the cut order are all the author’s own, written at `content/resources/research/respond/` on the website. This template operationalizes them and adds nothing to them.
- The scope sentences quoted in section 1 are from `sp2027.ieee-security.org/cfpapers.html`, the call for papers of the IEEE Symposium on Security and Privacy for 2027, read from the page.
- The survey of what response templates exist covers the Overleaf community templates tagged for responses, the Zenke Lab `LaTeX` template, van den Burg’s `LaTeXReviewReplyTemplate`, the `bamdome-rebuttal` Typst package, and Editage’s fillable form. SIGPLAN’s *Guidelines for the Program Chair of a SIGPLAN Event* is the venue-side statement.
- The percentages in section 2 are Table 3 of Kennard et al., *DISAPERE: A Dataset for Discourse Structure in Peer Review Discussions*, NAACL 2022.