

Filling the Response Template

Step seven, slot by slot, with the budget written down first

1 Where this sits

The response procedure has nine steps and it is on the website at `resources/research/respond/procedure.md`. Only step seven writes. This file is step seven opened up: which slot of `template_rebuttal.tex` gets filled, in what order, out of which input, and what goes wrong when the step is skipped.

The test for whether a step belongs here is the one the procedure already states. A step whose failure cannot be named is a habit rather than a procedure, and it can be dropped. Every step below names its failure.

2 What the search settled

Question 1 asks whether turning the advice into a template and a fill procedure is encouraged practice. The template half is answered in section 1 of `template_rebuttal.tex`: it is ordinary practice, and every published artifact is a journal response letter rather than anything a 750 word plain text box could hold.

The procedure half answers differently. Nothing published is a fill procedure. The response templates that exist are containers, a macro for a reviewer point and a macro for a reply, running in the order the reviews arrived. None of them supplies an order of work, a word budget, or a test for whether a filled one is finished. The advice runs the other way round, all order and no container. Noble gives ten rules and no skeleton. Sap gives heading practice and no skeleton. The venue calls give a scope sentence and a word count and stop.

So the pair has to be built, and the two halves have to be built against each other. Filling a container nobody wrote a procedure for is where the whole of question 1 sits, and the answer is that it is not discouraged. It is simply not done.

2.1 The three numbers that decide the order

DISAPERE labels every sentence of 506 review and response pairs from ICLR 2019 and 2020. Three of its counts change how this template is filled.

Answering is most of the work. Answering a question is 32.76 percent of all response sentences, three times the next argumentative label. Everything else in a response competes for what is left over.

One request in four asks for a run. Requests are 19.82 percent of review sentences, and inside that, requests for experiments are 4.78 against 14.42 for explanations, clarifications, edits and typos. Most of what a reviewer asks for closes with words.

Two sentences in five of a review are not a comment. Structuring, social text and the uncategorized remainder come to 38.82 percent of review sentences. That is the size of the pile the coverage audit is expected to close on inspection, and it is the reason the audit is cheap enough to run.

3 Six inputs, and what breaks without each

The fill step does not start until all six exist. They are products of steps 1 through 6, not of this one.

Input	What breaks without it
Venue mechanics	The draft is written for a formatted document and rewritten for a plain text box in the last hour, or the local count disagrees with the site's count at the deadline.
The extraction file	The draft argues against a remembered review, and it characterizes prior work instead of quoting it. Both are invisible to whoever wrote the draft and obvious to the reviewer who wrote the comment.
The reviewer cards	The response becomes three single-reviewer responses stapled together, and it gives up ground no comment asked for.
The clusters, each named by its root	Sections collect a topic and commit to none of it.
The classification	A politely worded comment naming a real defect receives a defense.
The gate decisions, signed off	A promise reaches the form that the author never agreed to make, and promises are enforceable.

4 The budget, written down first

The whole allowance is 750 words at ACSAC and at the IEEE Symposium on Security and Privacy. The ACSAC set came to seven clusters. That fixes the arithmetic before a word is drafted.

Seven headings at six words each is 42. One sentence of thanks at the top is another 15. What is left is 693 words over seven clusters, which is 99 each, which is four sentences of 25.

Four sentences. The four roles in the template are not a preference about style. They are what the budget allows: one answer, one ground, one consequence, one change. A fifth sentence anywhere is a sentence taken out of somewhere else.

The allowances are not equal, and they are written beside the headings. Initial score position bounds how far a review can move, so a review near the line is worth more words than one far from it. Deciding that with a number, before drafting, is the only point at which it can be decided on the evidence rather than on how the drafting happened to go.

The scaffolding has a measured size and it gets cut on purpose. At ICLR lengths, structuring, summary, social text, follow-up questions and the uncategorized remainder are 36.54

percent of response sentences. Held at that rate over 750 words, scaffolding would take 274 of them. Here it is seven headings and one sentence of thanks. No summary paragraph, no preview of what each cluster will do, no closing. Fifty-seven words, not 274.

5 The fill order

Eight passes, each across the whole document rather than down one cluster at a time. Free-writing, outlining, sketching, formalizing and drafting are five different activities and should not be attempted at once. The same holds for the six things a cluster has to do.

1. **Number the clusters.** Write the word allowance beside each heading, weighted toward the reviews that can move.

Skipped: drafting in order spends the budget in order. The first cluster comes out at full depth, the last three come out compressed, and the deepest answer lands wherever drafting started instead of where the movement is.

2. **Title every cluster, then stop.** Fill every `\rcluster` before any body exists. Kind first, subject second. The submitted v2 of this package is the worked set: a misunderstanding about sample efficiency, a clarification of the benefit of opponent modeling, a consideration of the ablation design, the motivation and hidden assumptions of the heuristics, the misleading terminology about simulation and emulation, the editorial changes, the potential experiments.

Skipped: a title naming only a topic collects everything on that topic and commits the section to none of it, so three questions get answered at a third of the depth each. That is inadequate explanation, the largest negative in the validated feature set at 0.31.

3. **Paste every echo.** Fill every `\recho` with the reviewer's own characters, copied out of the extraction file, never retyped from memory. A phrase longer than about ten words is not the distinctive part of it.

Skipped: a correct sentence the reader cannot attach to their own question is a true sentence answering nothing, and in a document read by three people with different concerns, nobody does the work of matching it.

4. **Run the gate before any change slot holds text.** Every request takes the three questions in order and comes back signed off, and only then do `\rchange`, `\rcommit` and `\roffer` get filled. Three requests in four ask for an explanation, a clarification, an edit or a typo fix, and those close inside the response.

Skipped: a commitment nobody agreed to make is discovered at export, with no time left to withdraw it. A program committee can reject a camera ready that fails to deliver what the response promised.

5. **Answer every cluster before grounding any of them.** Fill every `\ranswer` slot in one pass, direct and as asked, graded to the size of the error: yes, no, not quite, almost.

Skipped: grounding a cluster whose answer is still unwritten produces a paragraph of background with the answer at the bottom, or with no answer in it. A reviewer scanning for whether their point was addressed should not have to read to the end to find out.

6. **Ground, then consequence.** `\rground` carries the fact, the number, the equation or the passage. `\rconseq` carries what follows, formally and for whoever would use the system. A number goes in only where it changes a decision.

Skipped: an explanation that stops at the mechanism leaves the reviewer to work out the implication, and the implication is the part they asked about. A number with no denominator and no use is read as evidence and displaces evidence that would have done work.

7. **Audit against the reviews, not against the draft.** List every reviewer sentence that no cluster quotes, then sort the list into summary, praise and gap. Most of it closes on inspection. What survives gets a clause inside an existing cluster, or the one coverage sentence saying the concerns were read, are agreed with, and will be addressed in a revised version.

Skipped: auditing the draft against itself confirms only that the draft says what it says. Silence on a reviewer's question, however minor, is what the venues and the evidence both read as evasion.

8. **Export, count, hold.** Generate the plain text from this source rather than keeping a second copy by hand. Check that no macro, marker, identifier, link or gate survived. Count on the venue's own counter and treat it as authoritative. Over the limit, cut restatement, then spectator voice, then pointing, then the cluster answering the least movable review. Then it goes to the author, and the author submits.

Skipped: a second copy kept by hand drifts from the first, and the drift is found at the form, in the last hour, with no way to tell which version is the real one.

6 When the venue says nothing

Two silences, and they take opposite branches.

Silence about shape. No venue says how a response should be organized. The strongest statement from a venue body is SIGPLAN's guidance to its chairs, which encourages holding a response phase and stops there. That silence is a free choice, and the choice goes to the reader who is skimming. What reaches the decision is a chair's synthesis of whether the comments were met, so structure a chair can skim beats prose that repays a close read.

Silence about content. ACSAC does not say whether a new number may appear. The IEEE Symposium on Security and Privacy does say, and forbids new results outright in the non-interactive rebuttal. Where the venue is silent and its neighbor is explicit, the explicit rule is the safe reading and following it costs almost nothing. Security custom fills the ACSAC silence the same way: a small new number answering a reviewer's question is normal, a new topic is not.

Write down which branch was taken and why, in the package, so the next submission to the same venue does not argue it a second time.

7 Exit tests

Ten checks over the filled source. All of them are mechanical, which is the point. A check run as an intention is not run.

- Every `\rcluster` names a kind before its subject.

- Every cluster has exactly one `\recho`, and it sits in the first clause rather than in a sentence of its own.
- Every cluster has an `\ranswer`, and nothing precedes it but the echo.
- Every `\rcommit` names a configuration and the table it lands in.
- Every `\roffer` is numbered and says what it would show.
- No `\rchange` carries work that did not come back through the gate.
- No sentence is left with no reviewer comment behind it.
- Every reviewer question is either quoted inside a cluster or on the audit list with a decision written beside it.
- The venue’s own counter reads under the limit.
- The exported plain text contains none of the macros, and none of the identifiers.

8 Sources

- The nine steps, the gate, the classification, the clustering rule, the four sentence roles, the phrasing shelf and the cut order are the author’s own, at `content/resources/research/respond/` on the website.
- The cluster titles in step 2 are the section headings of `v2.tex`, the submitted response in this package.
- The sentence counts are Tables 2 and 3 of Kennard et al., *DISAPERE: A Dataset for Discourse Structure in Peer Review Discussions*, NAACL 2022.
- The odds ratio 0.31 for inadequate explanation is from *Rebuttals Move Peer-Review Scores, but Initial-Review Structure Bounds the Movement*, arXiv 2606.22166, 2026.
- The 750 word limit and the ban on new results are from `sp2027.ieee-security.org/cfpapers.html` and the ACSAC calls for papers. The venue-side statement on response phases is SIGPLAN’s *Guidelines for the Program Chair of a SIGPLAN Event*.